Matched-start estimator audit

A controlled comparison of fitted evaluation and direct simulator rollout

01 Fix inputs

- same trained policies
- same 225 test starts
- same reward and horizon
- same undiscounted target

Controlled setup

02 Evaluate estimators

- FQE value estimate
- direct rollout return
- identical policy pairs
- collision-free RNG streams

Estimator outputs

03 Pair by seed

- same seed both policies
- compute per-seed contrast
- retain paired observations
- preserve seed dependence

Paired contrast

04 Quantify uncertainty

- seed-level t interval
- paired bootstrap
- permutation sensitivity
- leave-one-out diagnostics

Robust inference

AUDIT OUTPUT Compare estimator gaps, quantify the paired contrast, and test the robustness of the conclusion.